

Introduction to Object-Oriented Programming

Polymorphism

Christopher Simpkins

`chris.simpkins@gatech.edu`

Introduction to Object-Oriented Programming


Today we'll learn how to combine all the elements of object-oriented programming in the design of a program that handles a company payroll. Object-oriented programming requires three features:

- Data abstraction with classes (encapsulation)
- Inheritance
- Dynamic method binding

That last part, dynamic method binding, provides for subtype *polymorphism*, which we'll learn today.

Class Hierarchies

Class hierarchies depict the superclass-subclass relationships between families of related classes. Consider:



- Employee is the superclass of HourlyEmployee and SalariedEmployee
- Employee is more general than HourlyEmployee and SalariedEmployee, e.g., there at least as many Employees as either HourlyEmployees or SalariedEmployees
- HourlyEmployee and SalariedEmployee are richer than Employee because they extend Employee with additional features

A SalariedEmployee Class

Let's add SalariedEmployee to our class hierarchy. Here are the important pieces:

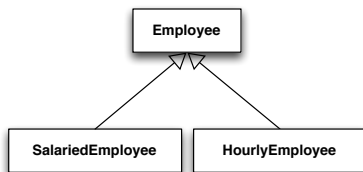
```
public final class SalariedEmployee3 extends Employee3 {

    private static final int MONTHS_PER_YEAR = 12;
    private final double annualSalary;

    public SalariedEmployee3(String aName, Date aHireDate,
                             double anAnnualSalary) {
        super(aName, aHireDate);
        disallowZeroesAndNegatives(anAnnualSalary);
        annualSalary = anAnnualSalary;
    }
    public double getAnnualSalary() {
        return annualSalary;
    }
    public double monthlyPay() {
        return annualSalary / MONTHS_PER_YEAR;
    }
    // ...
}
```

Our Employee Class Hierarchy

We now have all the classes in our hierarchy:



But our classes aren't well factored.

- `SalariedEmployee3` and `HourlyEmployee3` have duplicate copies of `disallowZeroesAndNegatives`
- `SalariedEmployee3` and `HourlyEmployee3` both have `monthlyPay` methods, but these methods are not polymorphic because they're not defined in `Employee3`

Let's refactor our `Employee` class hierarchy to give it a clean object-oriented design.

A Company Spec

Before we make `monthlyPay` polymorphic, we need an application to demonstrate why doing so is useful. Let's design a `Company` class with the following specs:

- A `Company`⁴ has exactly 9 employees (because we haven't learned about dynamically resized data structures yet)
- A company calculates its monthly payroll by adding up the monthly pay of each of its employees.
- A company can have any mix of hourly and salaried employees

That last bullet motivates the use of polymorphism.

Maintaining an Employee List

With our current class hierarchy, we need to maintain separate (partial) arrays of hourly and salaried employees. Because they're partial arrays we also need to keep track of how many of each type of employee we have.

```
public class Company {  
  
    private HourlyEmployee[] hourlyEmployees;  
    private int numHourlyEmployees = 10;  
    private SalariedEmployee[] salariedEmployees;  
    private int numSalariedEmployees = 10;  
  
    public Company() {  
        hourlyEmployees = new HourlyEmployee[numHourlyEmployees];  
        salariedEmployees = new SalariedEmployee[numSalariedEmployees];  
    }  
}
```

Calculating Payroll the Hard Way

With our employee lists, calculating payroll is accomplished with two loops:

```
public class Company { // hypothetical

    public double monthlyPayroll() {
        double payroll = 0.0;
        for (int i = 0; i < numHourlyEmployees; ++i) {
            payroll += hourlyEmployees[i].monthlyPay();
        }
        for (int i = 0; i < numSalariedEmployees; ++i) {
            payroll += salariedEmployees[i].monthlyPay();
        }
        return payroll;
    }
    // ..
}
```

Seems reasonable. But ...

- What if we want to add a third type of employee?

Calculating Payroll the Easy Way

We'd like to be able to calculate payroll with a single loop over all employees:

```
public class Company4 {  
  
    public double monthlyPayroll() {  
        double payroll = 0.0;  
        for (Employee employee: employees) {  
            payroll += employee.monthlyPay();  
        }  
        return payroll;  
    }  
    // ..  
}
```

Much cleaner and less error-prone (e.g., we don't have the book-keeping of two partial arrays). To be able to code like this we need to update the design of our `Employee` class hierarchy.

A More General Employee List

The first step is to store one array of `Employees`:

```
public class Company4 {  
    private Employee4[] employees;  
    public Company4() {  
        employees = ...;  
    }  
    public double monthlyPayroll() {  
        double payroll = 0.0;  
        for (int i = 0; i < employees.length; ++i) {  
            payroll += employees[i].monthlyPay();  
        }  
        return payroll;  
    }  
}
```

Much better. But it doesn't compile. Why?

```
$ javac Company.java  
Company.java:15: cannot find symbol  
symbol   : method monthlyPay()  
location: class Employee  
            payroll += employees[i].monthlyPay();
```

Abstract Classes

We need `Employee` to declare a `monthlyPay` method for subclasses to define. Since we don't have a general definition for `monthlyPay` suitable for `Employee`, `Employee` will need to be abstract.

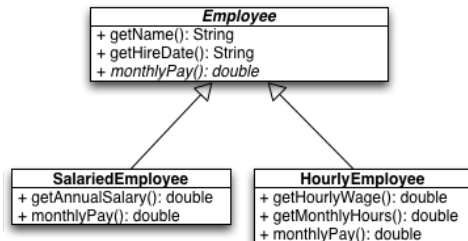
```
public abstract class Employee4 {  
    // ...  
    public abstract double monthlyPay();  
}
```

An abstract class

- cannot be instantiated,
- may contain zero or more abstract methods, and
- subclasses must either provide an implementation for abstract methods, or be declared `abstract` themselves.

This makes sense for our `Employee4` class. We don't ever want to instantiate `Employee4` objects. `Employee4` simply defines the common aspects of all employees, with subclasses filling in the details.

The Employee4 Class Hierarchy



- **Employee4** and its `monthlyPay` method are abstract.
- `monthlyPay` is polymorphic because it is overridden in subclasses.

Polymorphic Methods

```
public class Company4 {  
    private Employee4[] employees;  
    public double monthlyPayroll() {  
        double payroll = 0.0;  
        for (Employee4 employee: employees) {  
            payroll += employee.monthlyPay();  
        }  
        return payroll;  
    }  
}
```

- The static type of the elements of `employees` is `Employee4`
- The dynamic type can be any subclass of `Employee4`, in this case they are all `SalariedEmployee4` and `HourlyEmployee4`
- When a method is invoked on an object, the method of the dynamic (run-time) type is used, no matter what the static (compile-time) type is.
 - So though the static types of `employees` elements is `Employee`, the `monthlyPay` methods invoked on them are the ones defined in `SalariedEmployee4` and `HourlyEmployee4`.

Refactoring Duplicate Code in a Class Hierarchy

Recall the definition of `disallowZeroesAndNegatives`:

```
private void disallowZeroesAndNegatives(double ... args) {
    boolean shouldThrowException = false;
    String nonPositives = "";
    for (double arg: args) {
        if (arg <= 0.0) {
            shouldThrowException = true;
            nonPositives += arg + " ";
        }
    }
    if (shouldThrowException) {
        String msg = "Following arguments were <= 0: " + nonPositives;
        throw new IllegalArgumentException(msg);
    }
}
```

- This method is duplicated in `HourlyEmployee4` and `SalariedEmployee4`
- Let's move the definition of `disallowZeroesAndNegatives` into `Employee5` so it will be shared (rather than duplicated) in `SalariedEmployee5` and `HourlyEmployee5`.

protected Members

`private` members of a superclass are effectively invisible to subclasses. To make a member accessible to subclasses, use `protected`:

```
public abstract class Employee5 {  
    protected void disallowZeroesAndNegatives(double ... args) {  
        // ...  
    }  
    // ...  
}
```

`protected` members

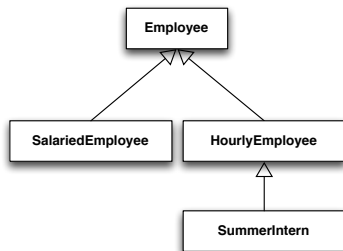
- are accessible to subclasses and other classes in the same package, and
- can be overridden in subclasses.

`protected` members provide encapsulation within a class hierarchy and package, `private` provides encapsulation within a single class.

Later we'll see a better way to re-use.

The Employee Class Hierarchy

Let's add a summer intern class to our Employee hierarchy.



- We can get the payroll for the current month by making use of the polymorphic `getMonthlyPay` method.
- What if we wanted to get the payroll for a particular month?

Let's overload `monthlyPay` so we can get the payroll for any month, not just the current month.

Enum Types

Enums are data types that have a predefined set of constant values ([JLS §8.9](#), [Java Enum Tutorial](#))

For example:

```
public enum Month {  
    JAN, FEB, MAR, APR, MAY, JUN, JUL, AUG, SEP, OCT, NOV, DEC  
}
```

defines an enum type called `Month` that can take on only one of the predefined constants `Month.JAN`, `Month.FEB`, ..., `Month.DEC`

- Enum types are a class.
- Java automatically defines convenience methods for enum types, like `valueOf(String)` and `values()` (See the [Enum API](#)).
- Because they define a class, enum types can include programmer-defined additional constructors and methods.

Ad-Hoc Polymorphism: Overloaded Methods

An overloaded method is a set of methods with the same names but different signatures (parameter lists)¹ (JLS §8.4.9).

Here's an overloaded `monthlyPay` for `SummerIntern6`, along with a helper method demonstrating the use of the `Month` enum:

```
public double monthlyPay() {
    Date today = new Date();
    Month thisMonth = Month.values()[today.getMonth()];
    return monthlyPay(thisMonth);
}

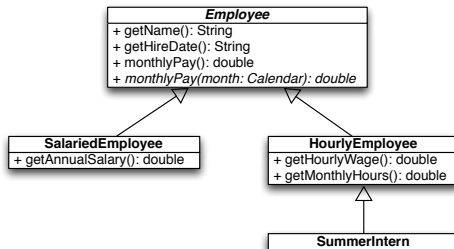
public double monthlyPay(Month month) {
    return isSummer(month) ? super.monthlyPay() : 0.0;
}

private boolean isSummer(Month month) {
    return month == Month.JUN
        || month == Month.JUL
        || month == Month.AUG;
}
```

■ In which classes should these methods be declared? Defined?

¹More precisely, two methods with the same name whose signatures are not

The Employee Class Hierarchy in UML



- Italicized names are abstract (e.g., *Employee* is an abstract class, *+ getMonthlyPay(month: Month)* is an abstract method).
- We've only shown public methods (denoted by the '+' symbols in front of their names).
- Each class has all the public methods in its superclasses, and possibly additional methods.
- SummerIntern *only specializes* HourlyEmployee, that is, it modifies some behavior of its superclass but does not add any additional behavior.

Forecasting Payroll

Now with our overloaded `monthlyPay` method we can forecast payroll:

```
Company6 c = new Company6();  
System.out.println("Monthly payroll this month: " +  
    c.monthlyPayroll());  
System.out.printf("Monthly payroll for May: %.2f%n",  
    c.monthlyPayroll(Month.MAY));  
System.out.printf("Monthly payroll for June: %.2f%n",  
    c.monthlyPayroll(Month.JUN));
```

Inheritance Hinders Re-use

Recall the `disallowZeroesAndNegatives` method that we refactored so that it's in the `Employee` class and inherited by subclasses:

```
public abstract class Employee6 {  
    protected void disallowZeroesAndNegatives(double ... args) {  
        // ...  
    }  
}
```

- There's nothing about this method that is specific to `Employees`
- `disallowZeroesAndNegatives` could be useful in other classes that are not part of the `Employee` class hierarchy.
- Since it's `protected`, it can't be used outside of the `Employee` class hierarchy or package.

In software engineering terms, we say that the code in `Employee` lacks *cohesion* - it has parts that aren't part of the *Employee* concept. Such a design hinders reuse.

Favor Composition over Inheritance

If we move these protected methods into a separate class, like [ValidationUtils.java](#)

```
public class ValidationUtils {  
  
    public static void disallowNullArguments(Object ... args) { ... }  
  
    public static void disallowZeroesAndNegatives(double ... args) {  
        ... }  
}
```

we can use them anywhere, e.g.,

```
public Employee(String aName, Date aHireDate) {  
    ValidationUtils.disallowNullArguments(aName, aHireDate);  
    name = aName;  
    hireDate = aHireDate;  
}
```

With this refactoring, we have our final versions of [Employee.java](#), [HourlyEmployee.java](#), and [SalariedEmployee.java](#)

Closing Thoughts on Polymorphism

We've now seen two kinds of polymorphism:

- Ad-hoc polymorphism (method overloading), and
- Subtype polymorphism (overriding methods in subtypes).

Subtype polymorphism is core feature of OOP. Polymorphism makes it possible to reuse *concepts* in a way that makes programs extensible without requiring rewriting existing code - this is the *open-closed principle*.

In the next block we'll see one more kind of polymorphism: type parameter polymorphism, or *parametric polymorphism*.

Object-oriented Design

With encapsulation, inheritance, and polymorphism we have all the language features we need to employ three important object-oriented design principles:

- **Single responsibility principle:** a class should only contain code related to the definition of the class
 - `Employee` classes contain employee-related code, validation code is in utility class
- **Open-closed principle:** open for extension, closed for modification
 - Can add new `Employee` subclasses without changing other classes in the `Employee` hierarchy or classes that use `Employees`, such as `Company`
- **Liskov substitution principle:** instances of subtypes should be substitutable wherever instances of supertypes are expected
 - A square is not a rectangle in an OO sense, but both are 2-D shapes

In CS 2340 you'll learn several more OO design principles and several patterns that employ them.